

Comps Paper

Alec Phillips

aphillips2@oxy.edu

Occidental College

1 Problem Context

Software testing is a core principle of computer science that is often overlooked or at least under-emphasized in computer science education. Testing is an integral part of any software project, and is a skill that must be developed like any other technique in software development. It takes considerable knowledge to be able to develop robust and efficient test cases that create confidence in one's code. Thus, this skill should be honed and developed along with other programming techniques. Software testing is also nuanced and there are numerous testing techniques that are important for students to learn about and have experience putting into practice. The importance of software testing serves as the motivation behind my comps project idea; I would like to create a web application that teaches software testing techniques and exposes introductory to intermediate level students to related concepts.

A web application that teaches software testing would be beneficial to computer science education in general because it could be added to a standard computer science curriculum to help students get exposure to the concept. My goal is to provide comprehensive materials on software testing as well as hands on exercises where learners can put the concepts into practice. Instructors could add this to their curricula, or self-learners could utilize it to gain exposure to software testing techniques.

2 Technical Background

This section will cover the relevant terminology that is critical for understanding the technical side of my project as well as the general goals of my application. These include general software testing techniques, as well as related topics such as test-driven development, error handling, edge case identification, and debugging. It is important that students view testing as an important aspect of software development.

2.1 Software Testing Techniques

2.1.1 Scopes of Testing

Testing is a nuanced aspect of software development, and there are several key techniques that are commonly employed for writing tests. These different techniques are often used together to comprehensively exercise code, as they work together to test a single application in different ways. Thus, it is not sufficient to know just one of these techniques; a competent developer should be familiar with all of them, in order to know what to utilize depending on their needs. Software testing can be broken up into four main categories depending on the scope and motivation of the test. These categories are unit testing, integration testing, end-to-end testing, and acceptance testing [11]. These are the four categories of testing that I plan to teach in my application, so it is important to elaborate on them in this section.

- Unit tests are those that have the narrowest scope. These are tests that apply to individual modules of code. Defining a specific unit of the code is up to the developer, but it is standard for unit tests to apply to specific functions, although it can also apply to an entire module.
- Integration testing checks that the individual units are interacting as expected. This means that the interfaces between the units as well as the general information flow between them should be exercised.
- System, or end-to-end testing evaluates the overall requirements of the entire application. System tests would generally involve providing inputs at the most general or external level, and evaluating that the outputs are as expected. For these tests to pass, all the internal processes of the aspect being tested must be functioning correctly.
- Acceptance testing is the final step, in which it is determined if the piece of software conforms to the general specification requirements, and also aligns with what the client or user is expecting [11].

2.1.2 Styles of Testing

In addition to the four categories of testing discussed above, there are also two styles of tests that serve different pur-

poses, and can each be applied to a either subset or all of the above categories. Once again, it is important to know both of these styles, as a thorough test suite should incorporate both. These two types are functional and structural testing. Functional is the more broad of the two, in the sense that any of the four categories of tests above can be functional in nature. Functional tests are also called 'black box' tests. These tests are ones that assume no knowledge of the underlying implementation of the code being tested, and only examine the expected behavior of the section of code being run. These are tests that could be written provided only with the API for a class, or the definition of a function, along with an understanding of the expected output. In contrast, structural tests, or 'white box' tests, take into account the implementation of the code. They are geared towards exercising specific sections of the implementation to make sure that the internal code is operating as expected. This is looking deeper than just expected outputs, but instead at the operations taking place within the code. Unit tests, integration tests, and system tests can all be structural, however acceptance tests cannot. This is because acceptance tests evaluate whether the system's behavior aligns with what the end user is expecting, and the user or client will (likely) not have any understanding of the internal implementation of the code [14].

2.1.3 Arrange, Act, Assert

Along with these types/styles of testing, my application includes information on testing best practices. These include the *arrange*, *act*, *assert* style of writing tests. This is the generally accepted format of structuring each test case.

1. The arrange step involves setting up any context necessary to run the test - this could include initializing objects and variables that will be involved in the test.
2. In the act step, function calls are made that exercise the behavior under test. Results of these calls are set to variables.
3. The variables resulting from the act step are then checked for correctness in the assert step. This involves some form of *assertion* statement that will see if the act step resulted in the expected values and throw an error if the result is not as expected.

2.1.4 Error Handling

Finally, the application exposes students to the concept of error handling in code. This is more tangentially related to testing but is still important to understand. It relates to software testing because the appropriate use of error handling makes code easier to test and debug. My application focuses on error handling to the extent of including a content section that goes into detail on how to raise errors in

javascript, how to catch and handle these errors, as well as how to test for errors being raised. It can be easily overlooked that test cases can be dedicated to checking for a particular type of error being raised in a certain situation, so it is important and relevant for students to learn about how to test for raised errors in this way.

2.2 Test-Driven Development

Initially, I had wanted my application to specifically have students practice the process of test-driven development. However, after beginning to develop coding exercises, it became clear that it is difficult to have users practice this development strategy within the confines of an application. This is because the process of test-driven development inverts the standard order of writing and then testing code, and can be abstract. I decided that it would be more engaging for students to be able to immediately run their test cases against a specific function. Thus, I did not include exercises specifically on test-driven development, but it is still explained in the Learn section of my application. Because TDD is increasingly being used in industry, students should have an understanding of the general process.

It is important to understand that test-driven development is not a testing strategy, such as the ones discussed in the prior section. Instead, it is a software development framework, meaning that it informs the entire software development process, not just the testing aspect [8].

Test-driven development fits within the Agile approach of software development [9]. Agile strategies tend to involve an iterative process that repeats until a project is complete, and is a common industry practice. This sits in contrast to older styles of development, such as Waterfall, which have more upfront design prior to coding. Test-driven development is also considered a practice of extreme programming, or XP, which takes basic development principles, such as testing, and emphasizes them to drive the entire development process [4].

Test-driven development lays out a detailed approach to developing software that is centered around allowing code testing to push forward the project, in the sense that test cases are written that in turn motivate the code that needs to be written.

Steps of TDD:

1. write test cases for the next unit of code being added to the project
2. write the code that the first test case is exercising (until it satisfies the test case)
3. refactor the code as needed (both function and test case)
4. re-run all existing test cases to check for regressions in the code base [2]

This process then repeats until the project is complete. Thus, this process informs the development of the entire project, not just the testing aspect.

3 Prior Work

This section explores some of the literature on incorporating software testing in to computer science education, as well as some of the areas that this can be improved. This research helped guide and motivate my project in terms of helping me determine what my application should focus on and the gaps that it is helping to fill.

3.1 Software Testing in Education

This section will discuss some of the strategies that have been used to teach software testing and test-driven development in the past. This will inform the approach that I take for teaching software testing in my application. I came across a significant amount of discussion on the reasons why software testing is undertaught, as well as why students find it to be an uninteresting concept to learn. Additionally, there was discussion of how test-driven development specifically could be integrated into computer science education.

The lack of emphasis on software testing in education stems from both students and instructors. Computer science courses as well as entire programs are already packed with material, so adding lessons or assignments focused solely on testing can be infeasible [6]. Additionally, students often find implementation of projects more exciting than testing their code, and testing can often appear tedious. Students can also be unmotivated to test their code, because they do not want to see it fail [3]. This is also something that I have noticed as a teaching assistant; students will sometimes write test cases that are very specific to their code, when more general test cases would not pass. From being a TA I have also noticed that, in general, students do not take writing test cases very seriously. Even when instructed to write tests, students in Data Structures will treat it as an afterthought, heavily prioritizing other aspects of their projects.

If the incentives were flipped, and students had more of a desire to write failing test cases, they would likely take the task more seriously. This can inform the way that I design my practice exercises in the application; instead of having students write code and then test it, I would provide them with code that they then need to exercise with test cases, and their performance on this activity will be dictated by how robust their tests are. One potential way this could work is by providing a function that fails on some edge case which the learner has to identify on their own and then write a test that causes the provided code to fail. This would incentivise the student to thoroughly understand both the goal of the

function, as well as the actual code, and would provide the desire to write a good test that fails the code.

There is also existing research on having students apply the concepts of test-driven development in order to learn computer science in general. This would make software testing a more integral part of every project that students take on, shaping the belief that testing is an important part of writing code. Some strategies for implementing this in an education setting were to make students fully responsible for demonstrating the correctness of their code [5]. This would mean that students would not be given any automated test results before submitting their code, and would instead be responsible for determining their code's correctness on their own. This may be unrealistic to implement in a classroom setting, because it would likely put an even greater strain on students, especially when computer science classes are already time consuming and challenging. However, I can base the exercises in my application around this idea, requiring students to write code along with test cases that must reach a certain amount of coverage of their code, or catch certain edge cases.

4 Ethical Considerations

An application geared towards education may initially appear to be ethical, however there are still potential moral issues. Any application or product geared towards users and has the potential to affect real world outcomes has the inherent possibility of affecting negative change. Additionally, there are a number of ethical concerns that arise from taking on the role of educator; for instance: what are the ethical and pedagogical obligations that an educator has to the students? There is also the topic of power, as this application's service will likely only affect those who already have access to computer science education, reinforcing the current power dynamics in terms of who has access to education and the ability to learn computer science. Additionally, it reinforces power dynamics existing around who has internet access. Despite the initially benign appearance of this project, these ethical considerations emphasize that there are still potential ethical issues present that need to be considered carefully.

4.1 Educator Obligations

Taking on the role of educator places one under an ethical burden, as they are entrusted to have expertise on the topic and the ability to adequately convey accurate information. The University of Michigan education department offers a simple overview of nine core ethical obligations that teachers should uphold. Several of the obligations that they bring up will be difficult for me to uphold within the context of

this project, specifically personal responsibility and competence.

The webpage defines personal responsibility as taking “responsibility for obstacles to student success and to work assiduously to ensure equitable access to learning opportunities” [12]. This poses an issue for my project because of the amount of time I will be able to dedicate to maintaining this project after it is completed and the impersonal nature of an educational application. If this application ends up actually getting deployed, it would be infeasible for me to be responsive to every user and their individual needs.

The article defines competence as “[developing] and continually [working] to improve instructional competence, and to strive to engage in professionally-justified teaching practice at all times” [12]. This is similarly infeasible because of the commitment that it would take to actively continue improving my instructional competence. Within the context of my project, it would require continued updating of the materials and content, as well as continuing to further my knowledge to better serve the users of the application, which are not responsibilities that I can confidently commit to at this time. These two examples represent the issues that come along with taking on the role of educator and building an educational site. In order to address this, it is important that, if this application is to be deployed, I only allow it to be active for as long as I can uphold these obligations.

4.2 Power Reinforcement

In addition to the general ethical issues of building an education based application, there is the problem of who will have access to this application, and how that reinforces existing power dynamics in the world of computer science and technology. Power and computer science are tightly intertwined; technology is nearly omnipresent in our world, and many of the most influential and powerful individuals are those who control large technology companies. One clear factor that differentiates power in computer science and technology is whether or not someone has access to the internet. Individuals who have easy access to the internet are much more easily able to interact with existing technologies and web applications, and also have much more access to learning how to interact with computers in general. The issue with this web application, and conveying educational material over the internet in general, is that it is only making information more accessible to those who can already easily access it. Even if my application successfully makes the information easier to learn or more enjoyable, it is not reaching a new audience or further distributing access to learning computer science. This then reinforces the existing power structures that exist in computer science.

There is already a large power difference between those who have access to the internet and those who do not. The

more effective my application is, the more I would be contributing to that power difference. The power differences caused by unequal internet access have been well documented, and play a role in global power dynamics as well. The advent of the internet has created easy access to information and communication, however this only applies to those who have access to it. Societies that do not have the luxury of widespread internet access do not get these same benefits, which increases the power differentials between these societies [7]. This emphasizes the importance of the relationship between internet access and power. Therefore, further empowering those with internet access by means of a web application that only those individuals can benefit from will inherently further these existing power dynamics.

Power differences as a result of internet access have been further exacerbated by the COVID-19 pandemic. This has been a prevalent issue in the academic sphere; students who have access to reliable internet can more easily and consistently access their courses, and do not have to worry about missing material because of something outside of their control like stable internet [10]. This emphasizes that, currently, those without internet access are experiencing even greater challenges as a result of the pandemic, so introducing more web based applications will further this divide.

These issues are important to consider, but my hope is that ultimately my application can provide more benefit to individuals and make computer science more accessible overall. Additionally, the way that my application is designed works to mitigate these concerns. Specifically, because the application can be fully contained within the user’s web browser and there are no network calls made during use, the application actually does not require an internet connection. If being used in the browser, a network connection is only required when the user initially requests the page. However, the application code can also be downloaded onto the device and can then be completely used with no internet connection. This makes the application more accessible and helps to combat some of the concerns above regarding the furthering of the power divide caused by internet access.

4.3 Ethics of Security

In my initial project design I had planned on having users log in to their own account to track their progress, which would require storing user information such as an identifier and password. This inherently brings up ethical concerns, because if the application is to be breached in some way, these users’ trusted information would no longer be safe. This places the burden of securing this information on the developer. However, once I switched to a serverless design and used the localStorage instead of a database to store progress, these concerns were eliminated. While I

was initially concerned about the limitations of building the application without a database, all the technical goals of the project were still able to be met, while also eliminating the concerns surrounding storing user information.

5 Methods

5.1 Implementation

The project is build fully using React.js and javascript, along with CSS and HTML. Originally, I intended to have a server-side API running that would process code submissions, as well as a database to store the excercises as well as user progress. This would require the application to have a login page, and the code evaluation would need to involve using subprocesses or something analogous in order to asynchronously run users' code submissions. Additionally, running user generated code on the server-side is inherently dangerous, because they gain access to the file system and OS of the server. In order to safely run their code, I would have needed to either restrict what the user is allowed to write in the code editor, or heavily verify their submissions. This would be very difficult and time consuming to do sufficiently, and would take away from the ability to focus on other aspects of the application like usability.

All of these factors lead me to end up choosing to build the application as a server-less web-page. This way, the user generated code is only running in their web-browser, and so they have no ability to break a remote server with their code. Additionally, this takes away the reliance on having a network connection to submit code, and there is no need to handle multiple submissions asynchronously, as would be necessary if a back-end were implemented. This also makes the application very fast, since it is independent of the speed of the network. One drawback of this design is that all coding exercises had to be in javascript, because that is the only language that can run in the web-browser, without the use of some sort of cross-compiler that would convert from another language into javascript. I had initially wanted to have the exercises be in python, however javascript is also a good language for students to get introduced to, especially if they are interested in going into front-end development.

One main drawback that the serverless design posed was the lack of a database. My main concern without a database was figuring out how to store the content descriptions and exercises, as well as storing user progress. However, I realized that another way to achieve the storing of user progress was through the browser's localStorage feature. This is a structure stored by the browser that each webpage can access through a simple get/set API. I leveraged this to store a list of exerciseIds corresponding to the exercises that the user has completed. Upon a correct exercise submission, the application updates this list, which is saved and loaded on

subsequent visits to the application. In order to get around storing other data that would have been put into a database, I stored the data in javascript objects. This simulates the way that the data would be stored in a non-relational database. I looked into serverless database options, but after learning that the localStorage option was available, I chose to go down that route because it seemed more feasible for the scale of the application.

Additionally, the lack of a server came along with several advantages for usability as well as ease of development. The general usability is improved because the lack of a back-end system means that the application is very quick and responsive, because there is no reliance on network connection speeds in sending information back and forth. Additionally, since all the data that the application needs in order to run is held in the browser upon initial loading, if network connection is subsequently lost, the application is still fully usable. This is advantageous because it makes the application more usable for individuals with less stable internet connections. In terms of ease of development, the lack of a server-side made the application much easier to deploy; I was able to deploy through github pages, which was a very simple process. This ease of deployment made it easier to have an initial deployment out earlier, meaning I was able to do a first round of user evaluations fairly early, allowing me more time to iterate on feedback and get input from users.

One additional aspect of the implementation that required outside research was getting code coverage for a user's test cases. From an initial search online, there were many tools that gather test coverage, but few that talk about how test coverage is actually determined. One technique that I located showed the process of injecting coverage flags into a function under test, which can be set to true when the execution enters the corresponding condition [1]. To implement this I simply declare an array of boolean flags, all set to false, with each flag corresponding to a particular branch of the code. When the function runs, within each branch, a different element of the array is set to true. Once all tests have been run, the total branch coverage is the percentage of flags that are set to true. This was a surprisingly simple way to determining coverage, and was effective for my purposes. There are downsides to evaluating simply on branch coverage, because there are more thorough methods for evaluating coverage, such as path coverage. I would have liked to extended to detecting path coverage with more time to work on the project, because it seems interesting algorithmically to detect, and would require users to write more thorough test cases, because there will necessarily be at least as many paths through a code block as there are branches.

5.2 Interface Design

My main motivation with the interface was to have a simple but still visually appealing design, focusing more on having a very intuitive page layout. My general design was inspired by CodingBat, which is one of my favorite online computer science education resources [13]. This site was created by a Stanford professor for his introductory computer science students. I find the overall experience of using the website to be intuitive and simple. The user interface is bare-bones and easy to navigate.

One difference that I wanted for my application was to have a more clear path through the material so that the content can build on itself. I wanted this for my application because I want the user to be working through the exercises in order so that they can draw on topics introduced earlier on in order to solve later problems. Additionally, since it is intended to be a first exploration of software testing as well as the javascript language, I believe that it makes sense for users to work through the problems from the beginning. In order to enforce this, I added the feature of unlocking exercises; at the beginning, the first exercise of each section is unlocked, and subsequent exercises become unlocked as the user completes the problems. I also believe that this makes the application more fun, because it gamify the user experience.

5.3 Content Design

As an educational application, a critical piece of my project was designing the content and programming exercises. I made sure to talk to many individuals about what would be most helpful to include and focus on in the application. The main takeaways were that it is important to include hands-on exercises as well as detailed examples, as opposed to simply having descriptions on types of testing. However, it was also emphasized that the materials should be concise to avoid being overwhelming. This was important to me as well, because I want this to be a concise but useful introduction to the nuances and types of software testing. One helpful piece of advice that I received on designing the exercises was to focus on common errors, such as off by one and improper input formats. By focusing on commonly occurring errors, users of the application will be able to gain familiarity with some of these and be better prepared to avoid them in the future.

In terms of the specific types of exercises included in the application, I had initially wanted to focus on implementing a combination of unit and system-style testing exercises. However, some aspects of development that I had overlooked took more time than expected, and I ended up focusing more on including a combination of edge-case identification and unit testing exercises, with minimal system

test exercises. This ended up being more feasible within the time constraints of the semester. However, I believe that the application still holds value as a teaching tool or for general student use. The edge-case identification exercises are still beneficial because it teaches students to reason about code correctness and consider the range of possible inputs and deeply understand the code presented. Additionally, these exercises prompt users to both identify an input that causes a function to break, and then actually debug the code. This is useful for gaining a deeper ability to reason about code, and offers an entry point into learning javascript for the user.

The test-writing exercises involve writing test cases that use assertion statements, with the goal of maximizing branch coverage on a given function. In the initial warm-up problems, the code being tested is provided, so the user can step through the function and more easily determine what inputs to provide to fully cover the branches. However, in subsequent exercises, a description of the behavior of the function is provided, but the actual code is obscured. This forces the user to consider how the function would be implemented in terms of the necessary control flow. To aid the user in this task, on a given submission attempt, they are given feedback on the percentage of branches that their code covered, as well as whether any of their assertion statements failed.

I ended up implementing one system-level test exercise, which involves the user writing a series of assertion tests on a larger scale API. This exercise is the most difficult of all the content and is only unlocked after a user has completed all the unit test exercises. This way they have the proper context to be able to complete the system test exercise. This exercise involves the user writing test cases against a mock course registration system. They are provided with the documentation for a CourseScheduler class with its member method stubs. Additionally, they are given the starting state of a 'database' which is actually implemented as a series of javascript objects. However, this simulates the experience of writing code that interacts with a back-end database. For this exercise, users are also trying to raise the overall branch coverage, and to do so must interact with all the methods in the CourseScheduler class and reason about what combinations of interactions will cause certain operations to occur. For instance, students need to exercise a prerequisite-checking function, but to do so they must enroll students in courses, increment the semester (so that students have then completed prerequisite courses), and then enroll the students in courses they are now qualified for.

5.4 Iterative Project Improvement

Since I was able to deploy an early version of the project, I had time to do two rounds of user evaluations, allowing me

to make iterative improvements to my project based on user feedback. In my first round of feedback, I used a survey that gauged the quality of the teaching content, the clarity of the exercise directions, the difficulty of several of the exercises, as well as overall user experience. From this first round of evaluations, I learned that the directions for the exercises could be made more clear, and that it would be helpful to include hints for some of the more difficult exercises. Additionally, it was noted that the feedback on coding exercises could be more detailed so that the user has more of a sense of what they did wrong. After this feedback, I added a hint area for each exercise, and improved the feedback reporting area to be more clear and visually appealing.

In addition to using a survey, I watched individuals interact with the application to get a sense of how clear it is to use and where individuals get confused. One main takeaway from this was that users sometimes found it difficult to know where to start with the coding exercises. The exercises are displayed in rows, and I intended for them to be attempted from top-left to bottom-right. However, I noticed that this was not always intuitive, and this informed my decision to add the feature of locking/unlocking exercises as the user progresses. This enforces the order and provides a clear visual of where to start.

This type of interactive testing also exposed some bugs in the code evaluation system, specifically with the system test exercise. One more advanced student tested the exercise and noticed that one specific series of interactions caused the exercise to break, so I was able to identify and fix the bug. Additionally, users identified multiple spelling errors in the application.

6 Evaluation Metrics

To formally evaluate my project, I chose to focus on three main criteria:

1. Teaching Quality (of written teaching content)
2. Exercise Quality (clarity of directions, difficulty)
3. Overall Learning Outcomes

As discussed above, these were evaluated in two stages, first one a larger scale with a survey that was sent out to members of the computer science community, and secondly with one on one user interviews where interactions with the application were observed.

6.1 Round 1 - User Survey

6.1.1 Survey Design

The survey was designed to target the three areas above for testing. Specifically, users were prompted to:

- read the content on unit testing
- read the descriptions of three exercises (one edge case identification, one debugging, one unit testing)
- attempt to complete the three problems within two to three minutes

This allowed me to gain a sense of how effective each of the evaluation criteria were met in an early version of the project. At this point in the implementation there were not yet system test exercises, so this aspect could not be evaluated. The problems that users were asked to complete were the rainfall problem for edge case and unit testing, and a fizzbuzz variation with an off by one error for debugging.

6.1.2 Participation

The survey was sent to ten upper division computer science students (students either majoring or minoring in computer science who have participated in an upper division course), as well as one section of twenty-eight introductory computer science students. Seven of the upper division students participated in the survey, and no introductory students participated. This was problematic because my application was intended to be geared specifically toward introductory students. However, the findings actually indicated that the more advanced students could still benefit from using the application, which makes me believe that a wider target audience could be appropriate.

6.2 Round 2 - One-on-one Evaluations

The second round of user evaluations involved asking individuals about their experience using the application in a one-on-one setting. I used two different formats for this.

- watching individual users interact with the application and observing with minimal interference
- asking individuals to use the application on their own (without observing them) and then discussing their experience

This was intended to gain insight into how users interact with the application when they have freedom to do so as they see fit. The decision to observe some users and not observe others was motivated by the fact that some users seemed to look for guidance when I was observing them using the application, which defeats the purpose of the user interacting with the application freely. However, it was still beneficial to get a direct visual of how individuals navigate the page. Additionally, the non-observational type allowed me to get feedback from individuals who were not directly present, such as peers who go to other schools.

With each of these tests, my evaluation consisted of an open ended discussion with the user about their experience,

as well as more targeted questions about their interactions. These questions focused on the three main evaluation criteria detailed above, so as to be consistent and work towards building a sense of the key goals of my project. However, these conversations also focused more on the overall usability of the application, and the clarity of the general application flow.

Overall these types of interviews were conducted with six students, three advanced and three introductory level.

7 Results and Discussion

7.1 Round One Evaluation

For the category of teaching quality, students indicated that, their prior knowledge of unit testing was 2.14 on a scale from one to five. After reading the learn section on unit testing, respondents rated their understanding at a 4 on the same one to five scale. Additionally, they rated the clarity of the section as 4.57 out of five. This was a satisfactory outcome, indicating that even for more advanced students, they felt that the section taught them something new. This is especially telling because the section they read was on unit testing, which I would have assumed to be the type of testing that advanced students would be most familiar with, however they still indicated positive learning outcomes.

For evaluating the exercise quality, students indicated the clarity of the exercise descriptions as well as whether they were able to solve the problem for each of the three types of problems that they attempted. Respondents reported average clarities of 3.86, 4, and 4.23 out of five for the edge case, debugging, and unit testing exercise directions respectively. These ratings indicate that users were generally able to determine what was expected for each type of exercise. However, in the written feedback area, some users indicated that the input format for edge case problems could be made more clear. In terms of difficulty, the solve rates were 71.43% (edge case), 85.71% (debugging), and 42.86% (unit test). These were all problems that I would have rated as easy to medium in terms of the problems that could have been chosen. Additionally, respondents were asked to only spend up to three minutes solving the problems. These results indicate that students had the most trouble solving the unit testing exercise within the time limit, which reinforces the general idea that students would benefit from such exercises/practice. However, this gave me a good baseline for understanding how my perception of problem difficulties translates to users' experiences.

In terms of overall learning outcomes, it becomes clear from the findings above that students seem to benefit from interacting with the application. This is emphasized by the fact that even upper division students felt that they gained knowledge from the written materials, and also were not

able to easily complete the unit testing exercise on the rainfall problem. Additionally, on a Likert scale from one to five, when asked if they would like this application to have been incorporated into their introductory computer science course:

- 14.29% gave a rating of 3
- 57.14% gave a rating of 4
- 28.57% gave a rating of 5

Overall the sentiment was positive and indicates that students would generally like to see this incorporated into their curriculum. Additionally, 100% of respondents indicated that they felt that they learned something through interacting with the application.

Despite the limited engagement with the survey, especially from introductory students, the findings were informative in providing a baseline of the applications effectiveness in relation to the three main evaluation criteria.

7.2 Round Two Evaluation

The one-on-one evaluations were particularly insightful for indicating how usable the application is, and whether it reached my goals for intuitiveness and clarity. Some main takeaways from watching individuals interact with the application were that introductory users tended to start with the learn section, while more advanced students (individuals who had taken at least one upper division computer science course) wanted to start with the exercises. In general, more advanced users also wanted to start with the harder exercises, instead of working their way up through the problems. In these cases, I would manually unlock the more advanced levels for them, but they tended to struggle because they were missing context that would have been explained if they had started from the easier problems. This reinforced my idea that it makes sense to still enforce the ordering of the problems even for more advanced users. In general, from watching the advanced students try exercises, they still seemed to find the problems somewhat challenging. Beginners also had difficulty, but were still able to solve the easier problems after reading the directions and contextual material. It was encouraging to see that students of all skill levels found the problems to be adequately difficult.

In terms of usability, testers indicated that the content sections were clear and they were not confused about where to navigate to in order to find the content that they wanted to look at. Additionally, users seemed to want to work from left to right through the content. There were still some points of confusion, specifically on the submission feedback for the coding problems. For the edge case identification problems, some users were confused because the goal of the problem is to provide an input that causes the function to execute incorrectly, but when they provide an

input that successfully does this, the feedback says 'PASS.' Users found this confusing because they thought that this meant that the function executed correctly, not that their input broke the function. I tried to fix this problem by using more visual cues that their solution passed, such as highlighting the submission green. Additionally, a green checkmark will appear on the problem name and the next exercise becomes unlocked. Given more time, I would further improve this by forcing the page to scroll back to the top or providing a 'next exercise' button to make it even more clear that the problem is complete. Additionally, for coding problems, the feedback indicates if there is an error in the user's code, but it does not say which line the error is on. Users complained that this made it difficult to debug their code. I tried to address this, but it seems like a limitation of the way the evaluation mechanism works; the user's code is embedded into an anonymous function that actually gets run in a virtual machine within the browser. I do not know all the specifics of why this is the case, but it seems to be the same for dynamically generated functions and any code run with the eval() function. I would also like to have tried other ways of addressing this if there was more time to work on the project.

7.3 Conclusion

Despite some of the shortcomings above, I believe that the project succeeded in meeting the goals that I set out to accomplish when designing the project. My main concern was that users would get a learning experience from using the application and be exposed to new content in a manageable and understandable way. From both sets of evaluations, it became clear that not only introductory, but also advanced students benefitted from interacting with the application. I believe that this indicates that the concept has more utility than I had initially thought. This also shows that more advanced students do not always learn software testing concepts independently even if they aren't taught these ideas in an introductory course. I believe that this reinforces the importance of introducing students to these concepts. In addition to the improvements discussed above, the main change that would be necessary if this application were to be used in a course is that more content would need to be added. It definitely proved to be more difficult than anticipated to come up with exercises of appropriate difficulty - often what seemed like a good idea for an exercise proved to be significantly too easy or difficult. It was specifically difficult to come up with unit test exercises, because the complexity or difficulty of an algorithm often did not translate to having a significant number of branches in the logic, meaning that it would only require a couple test cases to fully reach the coverage goal.

References

- [1] Baxter, Ira D. *Branch Coverage for Arbitrary Languages Made Easy*. Tech. rep. Austin, TX: Semantic Designs, Inc., 2001.
- [2] Bhat, Thirumalesh and Nagappan, Nachiappan. "Evaluating the Efficacy of Test-Driven Development: Industrial Case Studies". In: ISESE'06. Rio de Janeiro, Brazil: ACM, 2006.
- [3] Carrington, David. "Teaching Software Testing". In: ACSE'97. Melbourne, Australia: ACM, 1997.
- [4] Desai, Chetan, Janzen, David, and Savage, Kyle. "A Survey of Evidence for Test-Driven Development in Academia". In: *ACM SIGCSE Bulletin* 40.2 (2008), pp. 97–101.
- [5] Edwards, Stephen H. "Rethinking Computer Science Education from a Test-first Perspective". In: OOP-SLA. Anaheim, CA: ACM, 2003.
- [6] Edwards, Stephen H. "Teaching software testing: automatic grading meets test-first coding". In: OOP-SLA. Anaheim, CA: ACM, 2003.
- [7] Fang, Mei Lan et al. "Exploring Privilege in the Digital Divide: Implications for Theory, Policy, and Practice". In: *The Gerontologist* 59.1 (May 2018), e1–e15.
- [8] George, Boby and Williams, Laurie. "A structured experiment of test-driven development". In: *Information and Software Technology* 46 (2004), pp. 337–342.
- [9] Janzen, David and Saiedian, Hossein. *Test-Driven Development: Concepts, Taxonomy, and Future Direction*. Tech. rep. 0018-9162. Washington, DC: IEEE Computer Society, Sept. 2005.
- [10] Lai, John and Widmar, Nicole O. "Revisiting the Digital Divide in the COVID-19 Era". In: *Applied Economic Perspectives and Policy* 43.1 (2020), pp. 458–464.
- [11] Luo, Lu. *Software Testing Techniques: Technology Maturation and Research Strategy*. Tech. rep. Class Report for 17-939A. Pittsburgh, PA: Institute for Software Research International, Carnegie Mellon University, 2001.
- [12] Michigan School of Education, University of. *Ethical Obligations*. 2022. URL: <https://soe.umich.edu/academics-admissions/degrees/bachelors-certification/undergraduate-elementary-teacher-education/ethical-obligations>.
- [13] Parlante, Nick. *CodingBat: code practice*. 2017. URL: <https://codingbat.com/java>.

- [14] Sawant, Abhijit A., Bari, Pranit H., and Chawan, P.M. "Software Testing Techniques and Strategies". In: *IJERA* 2.3 (2012), pp. 980–986.